

Advanced Client Usage¶

Fork me on GitHub

Client Session¶

[ClientSession](#) is the heart and the main entry point for all client API operations.

Create the session first, use the instance for performing HTTP requests and initiating WebSocket connections.

The session contains a cookie storage and connection pool, thus cookies and connections are shared between HTTP requests sent by the same session.

Custom Request Headers¶

If you need to add HTTP headers to a request, pass them in a [dict](#) to the *headers* parameter.

For example, if you want to specify the content-type directly:

```
url = 'http://example.com/image'
payload = b'GIF89a\x01\x00\x01\x00\x00\xff\x00,\x00\x00'
          b'\x00\x00\x01\x00\x01\x00\x00\x02\x00;'
headers = {'content-type': 'image/gif'}

await session.post(url,
                   data=payload,
                   headers=headers)
```

You also can set default headers for all session requests:

```
headers={"Authorization": "Basic bG9naW46cGFzcw=="}
async with aiohttp.ClientSession(headers=headers) as session:
    async with session.get("http://httpbin.org/headers") as r:
        json_body = await r.json()
        assert json_body['headers']['Authorization'] == \
            'Basic bG9naW46cGFzcw=='
```

Typical use case is sending JSON body. You can specify content type directly as shown above, but it is more convenient to use special keyword *json*:

```
await session.post(url, json={'example': 'text'})
```

For *text/plain*

```
await session.post(url, data='Привет, Мир!')
```

Custom Cookies¶

To send your own cookies to the server, you can use the *cookies* parameter of [ClientSession](#) constructor:

```
url = 'http://httpbin.org/cookies'
cookies = {'cookies_are': 'working'}
async with ClientSession(cookies=cookies) as session:
    async with session.get(url) as resp:
        assert await resp.json() == {
            "cookies": {"cookies_are": "working"}}
```

Note

httpbin.org/cookies endpoint returns request cookies in JSON-encoded body. To access session cookies see [ClientSession.cookie_jar](#).

[ClientSession](#) may be used for sharing cookies between multiple requests:

```
async with aiohttp.ClientSession() as session:
    await session.get(
        'http://httpbin.org/cookies/set?my_cookie=my_value')
    filtered = session.cookie_jar.filter_cookies(
        'http://httpbin.org')
    assert filtered['my_cookie'].value == 'my_value'
    async with session.get('http://httpbin.org/cookies') as r:
        json_body = await r.json()
        assert json_body['cookies']['my_cookie'] == 'my_value'
```

Response Headers and Cookies¶

We can view the server's response [ClientResponse.headers](#) using a [CIMultiDictProxy](#):

```
assert resp.headers == {
    'ACCESS-CONTROL-ALLOW-ORIGIN': '*',
    'CONTENT-TYPE': 'application/json',
    'DATE': 'Tue, 15 Jul 2014 16:49:51 GMT',
    'SERVER': 'unicorn/18.0',
    'CONTENT-LENGTH': '331',
    'CONNECTION': 'keep-alive'}
```

The dictionary is special, though: it's made just for HTTP headers. According to [RFC 7230](#), HTTP Header names are case-insensitive. It also supports multiple values for the same key as HTTP protocol does.

So, we can access the headers using any capitalization we want:

```
assert resp.headers['Content-Type'] == 'application/json'
assert resp.headers.get('content-type') == 'application/json'
```

All headers are converted from binary data using UTF-8 with surrogateescape option. That works fine on most cases but sometimes unconverted data is needed if a server uses nonstandard encoding. While these headers are malformed from [RFC 7230](#) perspective they may be retrieved by using [ClientResponse.raw_headers](#) property:

```
assert resp.raw_headers == (
    (b'SERVER', b'nginx'),
    (b'DATE', b'Sat, 09 Jan 2016 20:28:40 GMT'),
    (b'CONTENT-TYPE', b'text/html; charset=utf-8'),
    (b'CONTENT-LENGTH', b'12150'),
    (b'CONNECTION', b'keep-alive'))
```

If a response contains some *HTTP Cookies*, you can quickly access them:

```
url = 'http://example.com/some/cookie/setting/url'
async with session.get(url) as resp:
    print(resp.cookies['example_cookie_name'])
```

Note

Response cookies contain only values, that were in Set-Cookie headers of the **last** request in redirection chain. To gather cookies between all redirection requests please use [aiohttp.ClientSession](#) object.

Redirection History¶

If a request was redirected, it is possible to view previous responses using the [history](#) attribute:

```

resp = await session.get('http://example.com/some/redirect/')
assert resp.status == 200
assert resp.url = URL('http://example.com/some/other/url/')
assert len(resp.history) == 1
assert resp.history[0].status == 301
assert resp.history[0].url = URL(
    'http://example.com/some/redirect/')

```

If no redirects occurred or `allow_redirects` is set to `False`, `history` will be an empty sequence.

Cookie Jar¶

Cookie Safety¶

By default [ClientSession](#) uses strict version of [aiohttp.CookieJar](#). [RFC 2109](#) explicitly forbids cookie accepting from URLs with IP address instead of DNS name (e.g. `http://127.0.0.1:80/cookie`).

It's good but sometimes for testing we need to enable support for such cookies. It should be done by passing `unsafe=True` to [aiohttp.CookieJar](#) constructor:

```

jar = aiohttp.CookieJar(unsafe=True)
session = aiohttp.ClientSession(cookie_jar=jar)

```

Dummy Cookie Jar¶

Sometimes cookie processing is not desirable. For this purpose it's possible to pass [aiohttp.DummyCookieJar](#) instance into client session:

```

jar = aiohttp.DummyCookieJar()
session = aiohttp.ClientSession(cookie_jar=jar)

```

Uploading pre-compressed data¶

To upload data that is already compressed before passing it to `aiohttp`, call the request function with the used compression algorithm name (usually `deflate` or `gzip`) as the value of the `Content-Encoding` header:

```

async def my_coroutine(session, headers, my_data):
    data = zlib.compress(my_data)
    headers = {'Content-Encoding': 'deflate'}
    async with session.post('http://httpbin.org/post',
                           data=data,
                           headers=headers) as pass

```

Disabling content type validation for JSON responses¶

The standard explicitly restricts JSON Content-Type HTTP header to `application/json` or any extended form, e.g. `application/vnd.custom-type+json`. Unfortunately, some servers send a wrong type, like `text/html`.

This can be worked around in two ways:

1. Pass the expected type explicitly (in this case checking will be strict, without the extended form support, so `custom/xxx+type` won't be accepted):

```
await resp.json(content_type='custom/type').
```

2. Disable the check entirely:

```
await resp.json(content_type=None).
```

Client Tracing

The execution flow of a specific request can be followed attaching listeners coroutines to the signals provided by the [TraceConfig](#) instance, this instance will be used as a parameter for the [ClientSession](#) constructor having as a result a client that triggers the different signals supported by the [TraceConfig](#). By default any instance of [ClientSession](#) class comes with the signals ability disabled. The following snippet shows how the start and the end signals of a request flow can be followed:

```
async def on_request_start(
    session, trace_config_ctx, params):
    print("Starting request")

async def on_request_end(session, trace_config_ctx, params):
    print("Ending request")

trace_config = aiohttp.TraceConfig()
trace_config.on_request_start.append(on_request_start)
trace_config.on_request_end.append(on_request_end)
async with aiohttp.ClientSession(
    trace_configs=[trace_config]) as client:
    client.get('http://example.com/some/redirect/')
```

The `trace_configs` is a list that can contain instances of [TraceConfig](#) class that allow run the signals handlers coming from different [TraceConfig](#) instances. The following example shows how two different [TraceConfig](#) that have a different nature are installed to perform their job in each signal handle:

```
from mylib.traceconfig import AuditRequest
from mylib.traceconfig import XRay

async with aiohttp.ClientSession(
    trace_configs=[AuditRequest(), XRay()]) as client:
    client.get('http://example.com/some/redirect/')
```

All signals take as a parameters first, the [ClientSession](#) instance used by the specific request related to that signals and second, a `SimpleNamespace` instance called `trace_config_ctx`. The `trace_config_ctx` object can be used to share the state through to the different signals that belong to the same request and to the same [TraceConfig](#) class, perhaps:

```
async def on_request_start(
    session, trace_config_ctx, params):
    trace_config_ctx.start = session.loop.time()

async def on_request_end(session, trace_config_ctx, params):
    elapsed = session.loop.time() - trace_config_ctx.start
    print("Request took {}".format(elapsed))
```

The `trace_config_ctx` param is by default a `SimpleNamespace` that is initialized at the beginning of the request flow. However, the factory used to create this object can be overwritten using the `trace_config_ctx_factory` constructor param of the [TraceConfig](#) class.

The `trace_request_ctx` param can given at the beginning of the request execution, accepted by all of the HTTP verbs, and will be passed as a keyword argument for the `trace_config_ctx_factory` factory. This param is useful to pass data that is only available at request time, perhaps:

```
async def on_request_start(
    session, trace_config_ctx, params):
    print(trace_config_ctx.trace_request_ctx)

session.get('http://example.com/some/redirect/',
            trace_request_ctx={'foo': 'bar'})
```

See also

[Tracing Reference](#) section for more information about the different signals supported.

Connectors¶

To tweak or change *transport* layer of requests you can pass a custom *connector* to [ClientSession](#) and family. For example:

```
conn = aiohttp.TCPConnector()  
session = aiohttp.ClientSession(connector=conn)
```

Note

By default *session* object takes the ownership of the connector, among other things closing the connections once the *session* is closed. If you are keen on share the same *connector* through different *session* instances you must give the *connector_owner* parameter as **False** for each *session* instance.

See also

[Connectors](#) section for more information about different connector types and configuration options.

Limiting connection pool size¶

To limit amount of simultaneously opened connections you can pass *limit* parameter to *connector*:

```
conn = aiohttp.TCPConnector(limit=30)
```

The example limits total amount of parallel connections to 30.

The default is 100.

If you explicitly want not to have limits, pass 0. For example:

```
conn = aiohttp.TCPConnector(limit=0)
```

To limit amount of simultaneously opened connection to the same endpoint ((host, port, is_ssl) triple) you can pass *limit_per_host* parameter to *connector*:

```
conn = aiohttp.TCPConnector(limit_per_host=30)
```

The example limits amount of parallel connections to the same to 30.

The default is 0 (no limit on per host bases).

Tuning the DNS cache¶

By default [TCPConnector](#) comes with the DNS cache table enabled, and resolutions will be cached by default for 10 seconds. This behavior can be changed either to change of the TTL for a resolution, as can be seen in the following example:

```
conn = aiohttp.TCPConnector(ttl_dns_cache=300)
```

or disabling the use of the DNS cache table, meaning that all requests will end up making a DNS resolution, as the following example shows:

```
conn = aiohttp.TCPConnector(use_dns_cache=False)
```

Resolving using custom nameservers¶

In order to specify the nameservers to when resolving the hostnames, [aiodns](#) is required:

```
from aiohttp.resolver import AsyncResolver

resolver = AsyncResolver(nameservers=["8.8.8.8", "8.8.4.4"])
conn = aiohttp.TCPConnector(resolver=resolver)
```

Unix domain sockets

If your HTTP server uses UNIX domain sockets you can use [UnixConnector](#):

```
conn = aiohttp.UnixConnector(path='/path/to/socket')
session = aiohttp.ClientSession(connector=conn)
```

SSL control for TCP sockets

By default *aiohttp* uses strict checks for HTTPS protocol. Certification checks can be relaxed by setting *ssl* to False:

```
r = await session.get('https://example.com', ssl=False)
```

If you need to setup custom ssl parameters (use own certification files for example) you can create a [ssl.SSLContext](#) instance and pass it into the proper [ClientSession](#) method:

```
sslcontext = ssl.create_default_context(
    cafile='/path/to/ca-bundle.crt')
r = await session.get('https://example.com', ssl=sslcontext)
```

If you need to verify *self-signed* certificates, you can do the same thing as the previous example, but add another call to [ssl.SSLContext.load_cert_chain\(\)](#) with the key pair:

```
sslcontext = ssl.create_default_context(
    cafile='/path/to/ca-bundle.crt')
sslcontext.load_cert_chain('/path/to/client/public/device.pem',
                          '/path/to/client/private/device.key')
r = await session.get('https://example.com', ssl=sslcontext)
```

There is explicit errors when ssl verification fails

[aiohttp.ClientConnectorSSLError](#):

```
try:
    await session.get('https://expired.badssl.com/')
except aiohttp.ClientConnectorSSLError as e:
    assert isinstance(e, ssl.SSLError)
```

[aiohttp.ClientConnectorCertificateError](#):

```
try:
    await session.get('https://wrong.host.badssl.com/')
except aiohttp.ClientConnectorCertificateError as e:
    assert isinstance(e, ssl.CertificateError)
```

If you need to skip both ssl related errors

[aiohttp.ClientSSLError](#):

```
try:
    await session.get('https://expired.badssl.com/')
except aiohttp.ClientSSLError as e:
    assert isinstance(e, ssl.SSLError)
```

```
try:
```

```

    await session.get('https://wrong.host.badssl.com/')
except aiohttp.ClientSSLSError as e:
    assert isinstance(e, ssl.CertificateError)

```

You may also verify certificates via *SHA256* fingerprint:

```

# Attempt to connect to https://www.python.org
# with a pin to a bogus certificate:
bad_fp = b'\0'*64
exc = None
try:
    r = await session.get('https://www.python.org',
                          ssl=aiohttp.Fingerprint(bad_fp))
except aiohttp.FingerprintMismatch as e:
    exc = e
assert exc is not None
assert exc.expected == bad_fp

# www.python.org cert's actual fingerprint
assert exc.got == b'...'

```

Note that this is the fingerprint of the DER-encoded certificate. If you have the certificate in PEM format, you can convert it to DER with e.g:

```
openssl x509 -in crt.pem -inform PEM -outform DER > crt.der
```

Note

Tip: to convert from a hexadecimal digest to a binary byte-string, you can use [binascii.unhexlify\(\)](#).

`ssl` parameter could be passed to [TCPConnector](#) as default, the value from [ClientSession.get\(\)](#) and others override default.

Proxy support

aiohttp supports HTTP/HTTPS proxies. You have to use *proxy* parameter:

```

async with aiohttp.ClientSession() as session:
    async with session.get("http://python.org",
                          proxy="http://proxy.com") as resp:
        print(resp.status)

```

It also supports proxy authorization:

```

async with aiohttp.ClientSession() as session:
    proxy_auth = aiohttp.BasicAuth('user', 'pass')
    async with session.get("http://python.org",
                          proxy="http://proxy.com",
                          proxy_auth=proxy_auth) as resp:
        print(resp.status)

```

Authentication credentials can be passed in proxy URL:

```

session.get("http://python.org",
            proxy="http://user:[email protected].")

```

Contrary to the requests library, it won't read environment variables by default. But you can do so by passing `trust_env=True` into [aiohttp.ClientSession](#) constructor for extracting proxy configuration from *HTTP_PROXY* or *HTTPS_PROXY* environment variables (both are case insensitive):

```

async with aiohttp.ClientSession(trust_env=True) as session:
    async with session.get("http://python.org") as resp:
        print(resp.status)

```

Proxy credentials are given from `~/.netrc` file if present (see [aiohttp.ClientSession](#) for more details).

Graceful Shutdown¶

When [ClientSession](#) closes at the end of an `async with block` (or through a direct [ClientSession.close\(\)](#) call), the underlying connection remains open due to `asyncio` internal details. In practice, the underlying connection will close after a short while. However, if the event loop is stopped before the underlying connection is closed, an `ResourceWarning: unclosed transport` warning is emitted (when warnings are enabled).

To avoid this situation, a small delay must be added before closing the event loop to allow any open underlying connections to close.

For a [ClientSession](#) without SSL, a simple zero-sleep (`await asyncio.sleep(0)`) will suffice:

```
async def read_website():
    async with aiohttp.ClientSession() as session:
        async with session.get('http://example.org/') as resp:
            await resp.read()

loop = asyncio.get_event_loop()
loop.run_until_complete(read_website())
# Zero-sleep to allow underlying connections to close
loop.run_until_complete(asyncio.sleep(0))
loop.close()
```

For a [ClientSession](#) with SSL, the application must wait a short duration before closing:

```
...
# Wait 250 ms for the underlying SSL connections to close
loop.run_until_complete(asyncio.sleep(0.250))
loop.close()
```

Note that the appropriate amount of time to wait will vary from application to application.

All if this will eventually become obsolete when the `asyncio` internals are changed so that `aiohttp` itself can wait on the underlying connection to close. Please follow issue [#1925](#) for the progress on this.



Async HTTP client/server for `asyncio` and Python



- build passing
- codecov 97%
- pypi package 3.4.4
- [gitter](#) [join chat](#)

Navigation

- [Client](#)
 - [Quickstart](#)
 - [Advanced Usage](#)
 - [Reference](#)
 - [Tracing Reference](#)
- [Server](#)
- [Utilities](#)

- [FAQ](#)
- [Miscellaneous](#)
- [Who uses aiohttp?](#)
- [Contributing](#)

Quick search

©2013-2018, Aiohttp contributors. | Powered by [Sphinx 1.7.7](#) | [Page source](#)